



Kaunas University of Technology
Faculty of Informatics

Report

JMH speed tests

P175B014 Data Structures

Author: Rasa Kundrotaitė, IFD-0/2
Teacher: Cenker Canbulut

Kaunas, 2021

Table of contents

What is JMH?	3
Research	3
Tests	3
Results	4
Sources	5

What is JMH?

Java microbenchmark harness (JMH) is a Java tool used to build, run and analyze nano/micro/mili/macro benchmarks written in Java and other languages targeting the JVM. Needless to say, the JMH tool does not guarantee that a benchmark is implemented correctly. According to its creators, there are multiple optimizations implemented on the JVM/OS or hardware side that make it challenging to measure the performance of a code.

Research

To start with, I have analyzed a few examples online on how JMH speed tests work ([source](#), [source](#)). After learning that it is possible to configure JMH in three ways (annotations, Java API, and Command-line), I have started testing the performance of my chosen methods:

Class BstSet: remove()	Class java.util.TreeSet<E>: remove()
------------------------	--------------------------------------

I have researched online their asymptotic complexities. [GeeksForGeeks](#) page has confirmed my hypothesis that for BstSet remove() method asymptotic complexity is $O(n)$. On [StackOverflow](#) I have found that TreeSet<E> remove() method has $O(\log(n))$ asymptotic complexity.

Tests

I have implemented two benchmark tests. The input data varied from 10000 to 100000. I have chosen to view the average time benchmark mode in μs .

These results can be seen from JMH speed tests:

Benchmark	(elementCount)	Mode	Cnt	Score	Error	Units
Benchmark.removeBstSet	10000	avgt	5	6.378 ±	0.408	us/op
Benchmark.removeBstSet	20000	avgt	5	13.952 ±	0.922	us/op
Benchmark.removeBstSet	40000	avgt	5	29.156 ±	2.499	us/op
Benchmark.removeBstSet	80000	avgt	5	57.983 ±	25.481	us/op
Benchmark.removeBstSet	100000	avgt	5	73.306 ±	7.340	us/op
Benchmark.removeTreeSet	10000	avgt	5	0.011 ±	0.001	us/op
Benchmark.removeTreeSet	20000	avgt	5	0.012 ±	0.004	us/op
Benchmark.removeTreeSet	40000	avgt	5	0.011 ±	0.001	us/op
Benchmark.removeTreeSet	80000	avgt	5	0.012 ±	0.001	us/op
Benchmark.removeTreeSet	100000	avgt	5	0.012 ±	0.001	us/op

The main parameters of the device used are below:

- Processor - AMD Ryzen 3 4300U with Radeon Graphics 2.70 GHz;
- Installed RAM - 8.00 GB (7.37 GB usable);
- System type - 64-bit operating system, x64-based processor;
- Pen and touch - not available.

Results

As we can see `TreeSet<E> remove()` method is being executed much quicker than `BstSet remove()` method. While `BstSet remove()` method's speed highly depends on the amount of input data (the more input data we provide, the slower the execution is), it is the opposite with `TreeSet<E> remove()` method, as its result for 100000 does not diverge much from 10000 input data. These results ideally represent the $O(n)$ asymptotic complexity of the `BstSet remove()` method and $O(\log(n))$ asymptotic complexity of `TreeSet<E> remove()` method. Overall, it would be wiser to use `TreeSet<E> remove()` method as it performs better than the `BstSet remove()` method.

Sources

1. <https://developpaper.com/how-to-benchmark-using-jmh-in-java/>
2. <https://www.awesome-testing.com/2019/05/performance-testing-benchmarking-java.html>
3. <https://www.geeksforgeeks.org/complexity-different-operations-binary-tree-binary-search-tree-avl-tree/>
4. <https://stackoverflow.com/questions/14379515/computational-complexity-of-treeset-methods-in-java>